

Part II: Generalized Linear Models

Load Packages

Again, we must load the packages that will be used in the first part of this workshop.

```
library(pastecs, quietly = TRUE)
library(lm.beta, quietly = TRUE)
library(lmtest, quietly = TRUE)
library(foreign, quietly = TRUE)
library(lattice, quietly = TRUE)
library(lme4, quietly = TRUE)
library(nlme, quietly = TRUE)
library(survival, quietly = TRUE)
library(dplyr, quietly = TRUE)
library(ggfortify, quietly = TRUE)
library(survminer, quietly = TRUE)
library(rms, quietly = TRUE)
library(MASS, quietly = TRUE)
library(pscl, quietly = TRUE)
```

Generalized linear models

A generalized linear model (GLM) has three:

- a random component with mean μ . Generally, the random component is the response variable Y_i .
- a systematic component, η_i , that relates the explanatory variables,

$$\eta_i = \sum_{j=1}^n \beta_j x_{ij}$$

- a link function that relates the mean of the random to the systematic component

$$g(\mu) = \eta_i$$

Logistic regression

Logistic regression is a GLM used the model binary (0 or 1) data. The response variable must be binary and is assumed to follow a bernoulli distribution.

That said, logistic regression has the following properties:

- a response binary variable, Y_i , that follows a bernoulli distribution with mean π_i .
- a systematic component, η_i , that relates the explanatory variables,

$$\eta_i = \sum_{j=1}^n \beta_j x_{ij}$$

- a link function that relates the mean of the random to the systematic component

$$\log\left(\frac{\pi_i}{1-\pi_i}\right) = \sum_{j=i}^n \beta_j x_{ij}.$$

$\log\left(\frac{\pi_i}{1-\pi_i}\right)$ is known as the log odds.

Data

Using the iris data, we create binary data. We add the column `Sepal.Width_binary` to `iris`. If the `Sepal.Width` is greater than the median then the associated value in `Sepal.Width_binary` is 1. Otherwise, `Sepal.Width_binary` is 0.

```
data <- iris
data$Sepal.Width_binary <- ifelse(data$Sepal.Width >= median(data$Sepal.Width), 1, 0)
```

Logistic Regression with only the constant term

Fitting only a constant term, the systematic component is

$$\eta_i = \beta_0$$

```
logit <- glm(Sepal.Width_binary ~ 1, data = data, family = "binomial")
summary(logit)
```

```
##
## Call:
## glm(formula = Sepal.Width_binary ~ 1, family = "binomial", data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.3911  -1.3911   0.9778   0.9778   0.9778
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   0.4895     0.1682    2.91  0.00361 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 199.22  on 149  degrees of freedom
## Residual deviance: 199.22  on 149  degrees of freedom
## AIC: 201.22
##
## Number of Fisher Scoring iterations: 4

p_avg <- mean(data$Sepal.Width_binary)
log_odds_avg <- log(p_avg/(1-p_avg))
print(log_odds_avg)

## [1] 0.4895482
```

Logistic Regression with Species

Fitting the species term, the systematic component is

$$\eta_i = 1 + \beta_2 X_{1i} + \beta_3 X_{2i}.$$

where

$$X_{1i} = \begin{cases} 1 & \text{if } i\text{th data point is versicolor} \\ 0 & \text{otherwise} \end{cases}, X_{2i} = \begin{cases} 1 & \text{if } i\text{th data point is virginica} \\ 0 & \text{otherwise} \end{cases}$$

```
logit <- glm(Sepal.Width_binary ~ as.factor(Species), data = data, family = "binomial")
summary(logit)
```

```
##
## Call:
## glm(formula = Sepal.Width_binary ~ as.factor(Species), family = "binomial",
##      data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5373  -0.8782   0.2857   1.0438   1.5096
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      3.1781     0.7215   4.405 1.06e-05 ***
## as.factor(Species)versicolor -3.9318     0.7826  -5.024 5.06e-07 ***
## as.factor(Species)virginica  -2.8553     0.7763  -3.678 0.000235 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 199.22  on 149  degrees of freedom
## Residual deviance: 147.51  on 147  degrees of freedom
## AIC: 153.51
##
## Number of Fisher Scoring iterations: 5
```

Let's compare the results to the average log odds of each Species group

```
log_odds_avg_fun <- function(data){
  p_avg <- mean(data)
  log_odds_avg <- log(p_avg/(1-p_avg))
  return(log_odds_avg)
}
```

```
tapply(data$Sepal.Width_binary,
       data$Species, log_odds_avg_fun)
```

```
##      setosa versicolor  virginica
## 3.1780538 -0.7537718  0.3227734
```

The intercept corresponds to the average log odds of setosa as we would expect. However, the other coefficients do not correspond to the average log odds of the other species. Why?

From the formula, $\eta_i = 1 + \beta_2 X_{2i} + \beta_3 X_{3i}$, the log odds of versicolor actually corresponds to $1 + \beta_2$. The log odds of virginica actually corresponds to $1 + \beta_3$.

```

coefficients<-unname(coef(logit))
print(c(coefficients[1],coefficients[1]+coefficients[2],
        coefficients[1]+coefficients[3]))

```

```
## [1] 3.1780537 -0.7537718 0.3227734
```

Logistic Regression with continuous variable

add discussion here

Logistic Regression with continuous variable, Sepal.Length

Fitting the species term, the systematic component is

$$\eta_i = \beta_3 X_{1i}.$$

where X_{1i} = Sepal.Length of the i th data point.

```

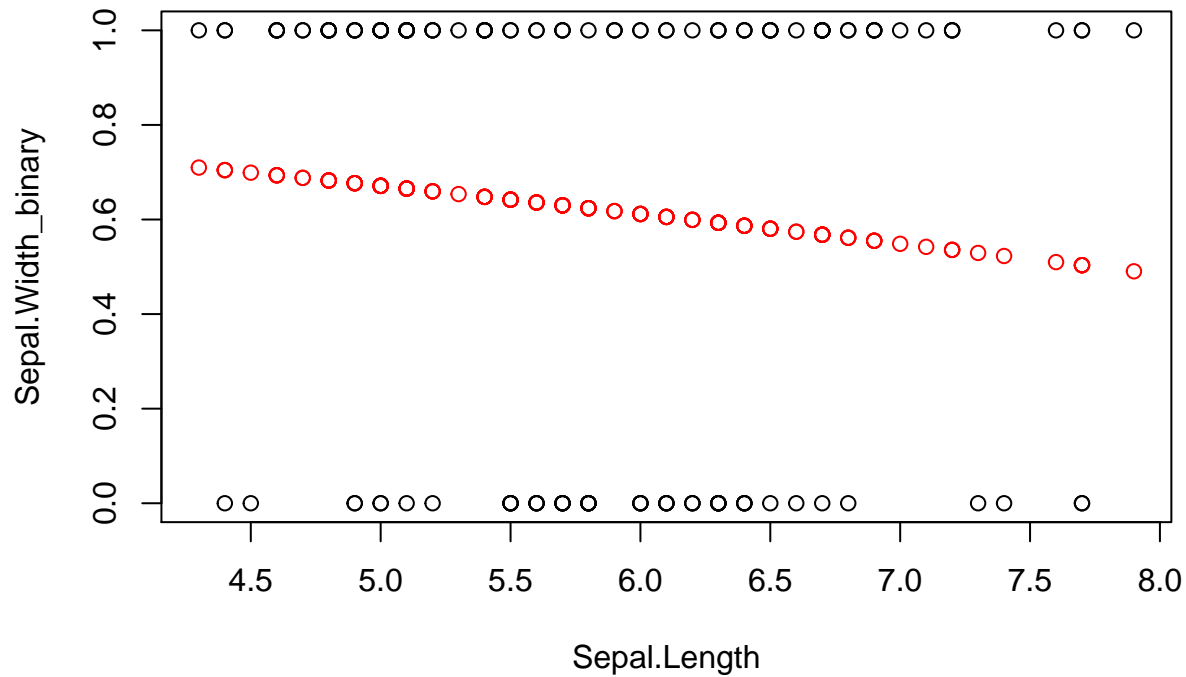
logit <- glm(Sepal.Width_binary ~ Sepal.Length,
             data = data, family = "binomial")
summary(logit)

##
## Call:
## glm(formula = Sepal.Width_binary ~ Sepal.Length, family = "binomial",
##      data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5614  -1.3524   0.8883   0.9890   1.1936
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    2.0088     1.2176   1.650   0.099 .
## Sepal.Length  -0.2591     0.2050  -1.264   0.206
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 199.22  on 149  degrees of freedom
## Residual deviance: 197.61  on 148  degrees of freedom
## AIC: 201.61
##
## Number of Fisher Scoring iterations: 4

plot(Sepal.Width_binary~Sepal.Length, data=data)
points(data$Sepal.Length[order(data$Sepal.Length)],
       logit$fitted[order(data$Sepal.Length)], col="red")
title(main="Data with Fitted Logistic Regression Line")

```

Data with Fitted Logistic Regression Line



Logistic Regression with Species and Sepal.Length

Fitting the species term, the systematic component is

$$\eta_i = 1 + \beta_2 X_{1i} + \beta_3 X_{2i} + \beta_3 X_{3i}.$$

where

$$X_{1i} = \begin{cases} 1 & \text{if } i\text{th data point is versicolor} \\ 0 & \text{otherwise} \end{cases}, \quad X_{2i} = \begin{cases} 1 & \text{if } i\text{th data point is virginica} \\ 0 & \text{otherwise} \end{cases}$$

and X_{3i} = Sepal.Length of the i th data point.

Fitting the logistic model accordingly,

```
logit <- glm(Sepal.Width_binary ~ Species + Sepal.Length,
             data = data, family = "binomial")
summary(logit)
```

```
##
## Call:
## glm(formula = Sepal.Width_binary ~ Species + Sepal.Length, family = "binomial",
##      data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2710  -0.7538   0.2472   0.7020   1.9477
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -4.7988     2.2981  -2.088 0.036784 *
```

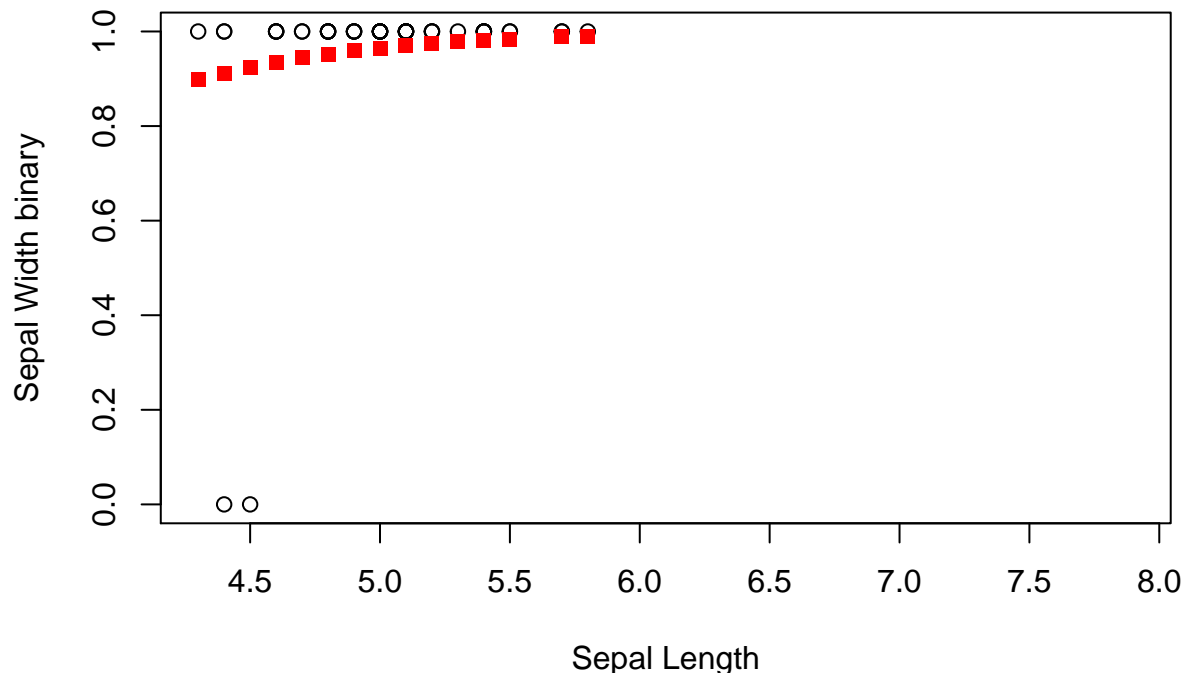
```
## Speciesversicolor -5.6936      0.9686 -5.878 4.16e-09 ***
## Speciesvirginica -5.4812      1.0879 -5.039 4.69e-07 ***
## Sepal.Length      1.6219      0.4510   3.596 0.000323 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 199.22  on 149  degrees of freedom
## Residual deviance: 131.27  on 146  degrees of freedom
## AIC: 139.27
##
## Number of Fisher Scoring iterations: 6
```

Plot the results for each species, we get that

```
plot(data[data$Species == "setosa", ]$Sepal.Length,
      data[data$Species == "setosa", ]$Sepal.Width_binary,
      xlim=as.matrix(range(data$Sepal.Length)),
      xlab = 'Sepal Length', ylab= 'Sepal Width binary',
      main= 'Scatter plot of sepal length vs sepal width')

points(data$Sepal.Length[data$Species == "setosa"],
        logit$fitted[data$Species == "setosa"], pch=15,
        col="red")
```

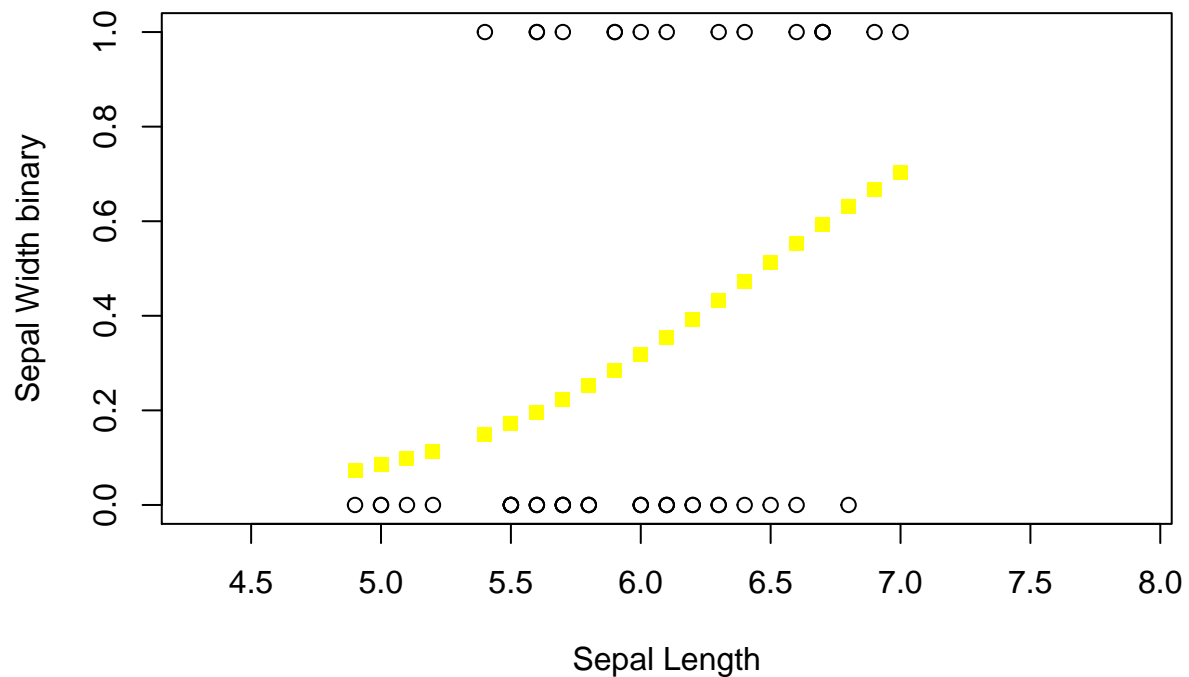
Scatter plot of sepal length vs sepal width



```
plot(data[data$Species == "versicolor", ]$Sepal.Length,
      data[data$Species == "versicolor", ]$Sepal.Width_binary,
      xlim=as.matrix(range(data$Sepal.Length)),
      xlab = 'Sepal Length', ylab= 'Sepal Width binary',
      main= 'Scatter plot of sepal length vs sepal width')
```

```
points(data$Sepal.Length[data$Species == "versicolor"],
       logit$fitted[data$Species == "versicolor"], pch=15,
       col="yellow")
```


Scatter plot of sepal length vs sepal width



```
plot(data[data$Species == "virginica", ]$Sepal.Length,
     data[data$Species == "virginica", ]$Sepal.Width_binary,
     xlim=as.matrix(range(data$Sepal.Length)),
     xlab = 'Sepal Length', ylab= 'Sepal Width binary',
     main= 'Scatter plot of sepal length vs sepal width')

points(data$Sepal.Length[data$Species == "virginica"],
       logit$fitted[data$Species == "virginica"], pch=15,
       col="blue")
```

Scatter plot of sepal length vs sepal width



Goodness of Fit

Deviance

For general linear models, we use *deviance* to compare two different models. Deviance is the difference in log likelihood of the models multiplied by 2.

Saturated Model

Let's consider a model in which each data point has its own mean and coefficients. This is called the saturated model. It basically replicates the data at hand.

Using deviance, we can compare our fitted model to a saturated model. If the fitted model behaves similarly to the saturated model, then the deviance can be well approximated by a chi-squared distribution with $m - n$ degrees of freedom. m is the number of data points and n is the number of coefficients in our fitted model.

This statistical property of the deviance allows us to perform a hypothesis test

H_0 : the fitted model is equivalent to the saturated model

H_a : the fitted model is not equivalent to the saturated model

`logit$deviance` is the deviance between the saturated model and the fitted model. `logit$df.residual` is equal to the number of observations minus the number of coefficients in the fitted model. Using this, we can calculate the p-value for the hypothesis test above.

```
p_value = pchisq(logit$deviance,
                  logit$df.residual, lower.tail = F)
print(p_value)
```



```
## [1] 0.8032738
```

Since the p value is greater than 0.05, we fail to reject the null hypothesis. (This is a good thing.)

Null Model

We can also use deviance to determine if our fitted model is better than the null model. The null model is a model with only a linear term. Like above, we can design a hypothesis test comparing the null model to the fitted model.

$$H_0 = \text{the fitted model is equivalent to the null model}$$
$$H_\alpha = \text{the fitted model is not equivalent to the null model}$$

In the limit of large data, it is known that the deviance follows a chi-squared distribution with parameter $n - 1$.

`logit$deviance` is the deviance between saturated model and fitted model. `logit$df.residual` is equal to number of observations minus the number of coefficients in the fitted model.

`logit$null.deviance` is the deviance between saturated model and the null model. `logit$df.null` is the number of observations minus 1.

Using this information, we can calculate the p value for the hypothesis test above.

```
p_value = pchisq(logit$null.deviance-logit$deviance,
                  logit$df.null-logit$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 1.173879e-14
```

Since the p value is less than one, we reject our null hypothesis. (This is a good thing.)

Anova

`anova` with argument `test="Chisq"` allows us to compare change in deviance after sequentially adding terms our model.

```
anova(logit,test="Chisq")
```

```
## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Sepal.Width_binary
##
## Terms added sequentially (first to last)
##
##
##           Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                149      199.22
## Species             2    51.709      147    147.51 5.910e-12 ***
## Sepal.Length        1    16.239      146    131.27 5.583e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Poisson General Linear Model

A poisson GLM is used to study *count* data (i.e. discrete numbers, $0, 1, 2, \dots$). *Count* data describes the number of events that occur within a given time frame.

insert plot of poisson distribution here

A poisson GLM is most useful when studying data in which the mean and variable are approximately equal. If they are not equal, the standard error of the model terms must be adjusted to account for the assumption violation.

Poisson Regression has the following properties:

- response count variables, Y_i , that follows a Poisson distribution with mean μ_i
- a systematic component, η_i , that relates the explanatory variables, $\eta_i = \sum_{j=1}^n \beta_j x_{ij}$
- a link function, $\log(\mu_i) = \sum_{j=1}^n \beta_j x_{ij}$

From Poisson regression, we learn the *mean* of each Y_i given the associated explanatory variables.

Data

We will consider the `bioChemists` data set in this section. This data set contains number of articles produced by PhD biochemistry student during the last 3 years of their PhD.

```
attach(bioChemists)
summary(bioChemists)
```

```
##      art      fem      mar      kid5
## Min.   : 0.000  Men :494  Single :309  Min.   :0.0000
## 1st Qu.: 0.000  Women:421  Married:606  1st Qu.:0.0000
## Median : 1.000                                Median :0.0000
## Mean   : 1.693                                Mean   :0.4951
## 3rd Qu.: 2.000                                3rd Qu.:1.0000
## Max.   :19.000                                Max.   :3.0000
##      phd      ment
## Min.   :0.755  Min.   : 0.000
## 1st Qu.:2.260  1st Qu.: 3.000
## Median :3.150  Median : 6.000
## Mean   :3.103  Mean   : 8.767
## 3rd Qu.:3.920  3rd Qu.:12.000
## Max.   :4.620  Max.   :77.000
```

The data set also contains demographic data associated with each student. data of the flower of certain plant species. The data set has five variables:

- *art* - number of articles produced by the student in the last 3 years of their PhD
- *fem* - gender
- *mar* - marital status
- *kid5* - number of children less than 5
- *phd* - prestige of PhD program
- *ment* - number of articles of the mentor in the last 3 years

```
sapply(bioChemists, class)
```

```
##      art      fem      mar      kid5      phd      ment
## "integer" "factor" "factor" "numeric" "numeric" "integer"
```

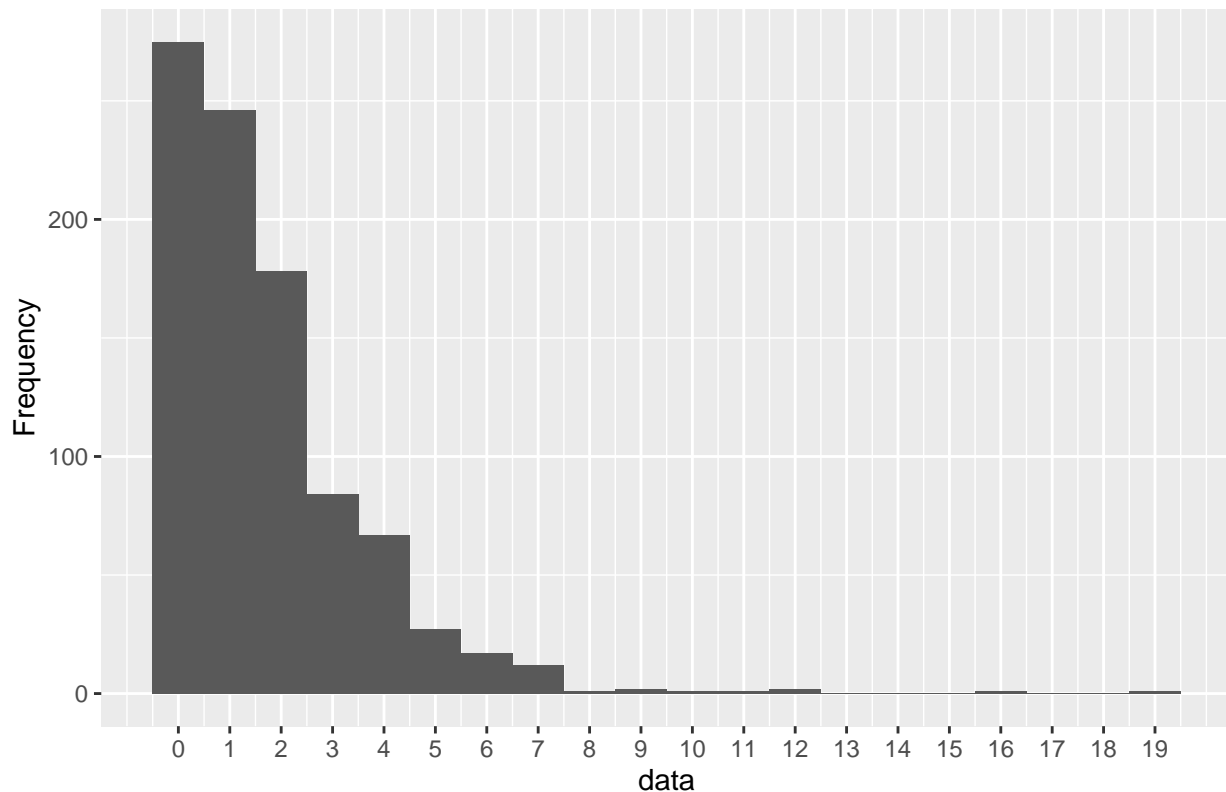
I convert `bioChemists$kid5` from numeric to factor. This will be used later.

```
bioChemists$kid5 <- factor(bioChemists$kid5,
                           levels= unique(bioChemists$kid5),
                           labels= unique(bioChemists$kid5))
```

Plotting the bar graph of `bioChemists$art`, we can see than the data looks Poisson-like since there is large number of observations at 0.

```
ggplot(bioChemists,aes(x=bioChemists$art))+
  geom_histogram(binwidth = 1, center = 1) +
  scale_x_continuous(breaks=seq(0,max(bioChemists$art), by = 1))+
  ylab("Frequency")+ xlab("data")+
  ggtitle("Histogram plot of the number of articles published by biochemist phd students in last 3 year")
```

Histogram plot of the number of articles published by biochemist phd stude



We can “quantify” the Poisson-ness by analyzing the mean and variance of the data.

```
mean(bioChemists$art)
```

```
## [1] 1.692896
```

```
var(bioChemists$art)
```

```
## [1] 3.709742
```

Although mean and variance are not equal, we will still fit it to Poisson distribution.

Possion Regression with constant term

To model only the constant term, I use the formula `art ~ 1`. This formula is equivalent to

$$\log \mu_i = \beta_0.$$

```
poisson_model = glm(art ~ 1, family=poisson(link=log), data=bioChemists)
summary(poisson_model)

##
## Call:
## glm(formula = art ~ 1, family = poisson(link = log), data = bioChemists)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8401  -1.8401  -0.5770   0.2294   7.5677
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.52644    0.02541   20.72  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 1817.4  on 914  degrees of freedom
## Residual deviance: 1817.4  on 914  degrees of freedom
## AIC: 3487.1
##
## Number of Fisher Scoring iterations: 5
```

Note that the constant term is the log mean number of counts.

```
print(coef(poisson_model))

## (Intercept)
##      0.5264408

print(log(mean(bioChemists$art)))

## [1] 0.5264408
```

Goodness of fit

Saturated model

We can again compare the current model to the saturated model (best possible fit).

```
p_value = pchisq(poisson_model$deviance,
                 poisson_model$df.residual, lower.tail = F)
print(p_value)

## [1] 3.304511e-62
```

Since our p value is less than 0.05, we reject the null hypothesis. The models are not equivalent.

Null model

We can also compare the current model to the null model (worst possible fit).

```
p_value = pchisq(poisson_model$null.deviance-poisson_model$deviance,
                 poisson_model$df.null-poisson_model$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 1
```

We fail to reject the null hypothesis. This makes sense since the models are literally the same thing.

Poisson Regression with marital status covariate

To model the marital status covariate, I use the formula `art ~ 1+mar`. This formula is equivalent to

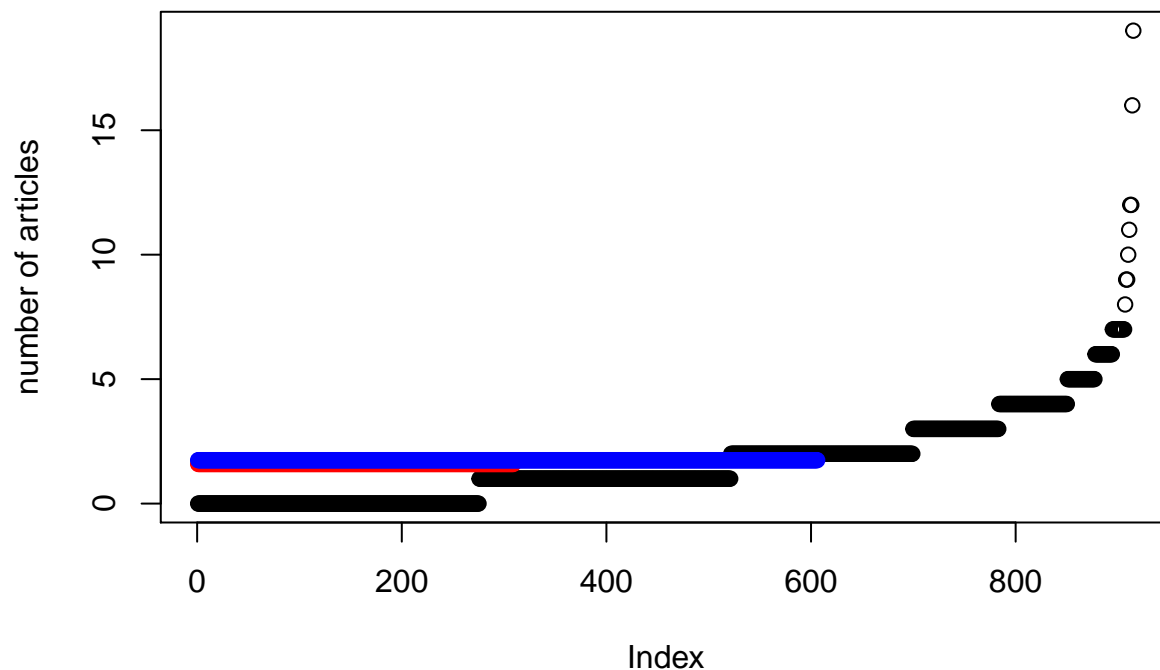
$$\log \mu_i = \beta_0 + \beta_1 X_{1i}$$

where

$$X_{1i} = \begin{cases} 1 & \text{if mar = Married} \\ 0 & \text{otherwise} \end{cases}.$$

```
poisson_model = glm(art~1+mar , family=poisson(link=log),data=bioChemists)
summary(poisson_model)

##
## Call:
## glm(formula = art ~ 1 + mar, family = poisson(link = log), data = bioChemists)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.8677  -1.7845  -0.5042   0.3107   7.4992
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.46514    0.04508  10.317  <2e-16 ***
## marMarried   0.09117    0.05458   1.671   0.0948 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 1817.4  on 914  degrees of freedom
## Residual deviance: 1814.6  on 913  degrees of freedom
## AIC: 3486.3
##
## Number of Fisher Scoring iterations: 5
plot(bioChemists$art,ylab='number of articles',xlab = 'Index')
points(poisson_model$fitted[bioChemists$mar=='Single'],col="red")
points(poisson_model$fitted[bioChemists$mar=='Married'],col="blue")
```



Graphically, we can see that marital status is not a good indicator of the number of articles published.

Goodness of fit

Saturated model

We can again compare the current model to the saturated model (best possible fit).

```
p_value = pchisq(poisson_model$deviance,
                 poisson_model$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 4.731233e-62
```

Since our p value is less than 0.05, we reject the null hypothesis. The models are not equivalent and our model is a bad fit.

Null model

We can also compare the current model to the null model (worst possible fit).

```
p_value = pchisq(poisson_model$null.deviance-logit$deviance,
                 poisson_model$df.null-logit$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 1.016236e-70
```

Since our p value is less than 0.05, we reject the null hypothesis. The models are not equivalent. Though our current model does not capture much deviance, the current model captures much more variance than the null model.

```
anova(poisson_model, test="Chisq")
```

```
## Analysis of Deviance Table
##
## Model: poisson, link: log
##
## Response: art
##
## Terms added sequentially (first to last)
##
##
##      Df Deviance Resid. Df Resid. Dev Pr(>Chi)
## NULL                914      1817.4
## mar    1    2.8211      913      1814.6 0.09304 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Possion Regression with martial status and children covariate

To model the martial status and children as covariates, I use the formula `art ~ 1+mar + kid5`. This formula is equivalent to

$$\log \mu_i = \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i} + \beta_3 X_{3i} + \beta_4 X_{4i}$$

where

$$X_{1i} = \begin{cases} 1 & \text{if the } i\text{th data point is married} \\ 0 & \text{otherwise} \end{cases},$$

$$X_{2i} = \begin{cases} 1 & \text{if the number of children of } i\text{th data point is 1} \\ 0 & \text{otherwise} \end{cases},$$

$$X_{3i} = \begin{cases} 1 & \text{if the number of children of } i\text{th data point is 2} \\ 0 & \text{otherwise} \end{cases}$$

$$\text{and } X_{4i} = \begin{cases} 1 & \text{if the number of children of } i\text{th data point is 3} \\ 0 & \text{otherwise} \end{cases}.$$

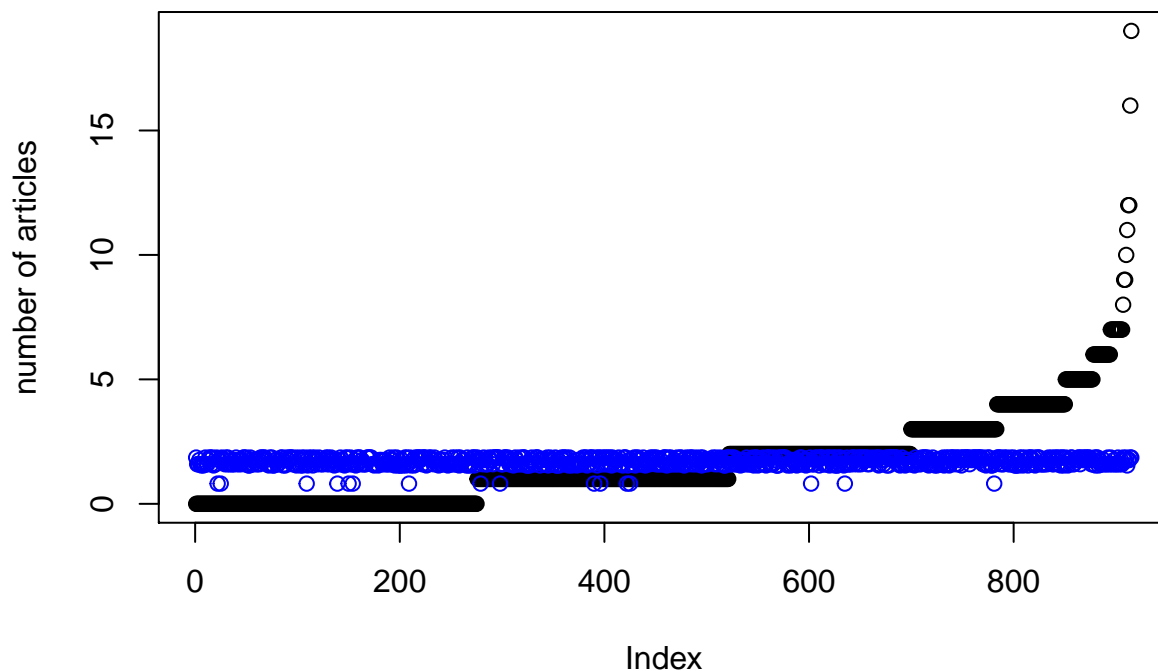
```
poisson_model = glm(art ~ 1 + kid5 + mar,
                    family=poisson(link=log),data=bioChemists)
summary(poisson_model)
```

```
##
## Call:
## glm(formula = art ~ 1 + kid5 + mar, family = poisson(link = log),
##      data = bioChemists)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9280  -1.7845  -0.5042   0.3518   7.3520
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   0.46514    0.04508  10.317  <2e-16 ***
## kid51         -0.05510    0.06907  -0.798   0.4250
```

```
## kid52      -0.18620    0.08960   -2.078    0.0377 *
## kid53      -0.82747    0.28067   -2.948    0.0032 **
## marMarried  0.15470    0.06235    2.481    0.0131 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
## Null deviance: 1817.4  on 914  degrees of freedom
## Residual deviance: 1799.9  on 910  degrees of freedom
## AIC: 3477.7
##
## Number of Fisher Scoring iterations: 5
```


```
plot(bioChemists$art,ylab='number of articles',xlab = 'Index')
single_bioChemists = bioChemists[bioChemists$mar=='Single',]

points(poisson_model$fitted[single_bioChemists$kid5== 0],col="blue",pch=1)
points(poisson_model$fitted[single_bioChemists$kid5== 1],col="yellow",pch=1)
points(poisson_model$fitted[single_bioChemists$kid5== 2],col="red",pch=1)
points(poisson_model$fitted[single_bioChemists$kid5== 3],col="green",pch=1)
```



```
plot(bioChemists$art,ylab='number of articles',xlab = 'Index')
mar_bioChemists = bioChemists[bioChemists$mar=='Married',]

points(poisson_model$fitted[mar_bioChemists$kid5== 0],col="blue",pch=2)
points(poisson_model$fitted[mar_bioChemists$kid5== 1],col="yellow",pch=2)
points(poisson_model$fitted[mar_bioChemists$kid5== 2],col="red",pch=2)
points(poisson_model$fitted[mar_bioChemists$kid5== 3],col="green",pch=2)
```

Graphically, we can see that marital status and number of children is not a good indicator of the number of articles published.

Goodness of Fit

Saturated model

We can again compare the current model to the saturated model (best possible fit).

```
p_value = pchisq(poisson_model$deviance,
                  poisson_model$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 6.462874e-61
```

Since our p value is less than 0.05, we reject the null hypothesis. The models are not equivalent.

Null model

We can also compare the current model to the null model (worst possible fit).

```
p_value = pchisq(poisson_model$null.deviance - poisson_model$deviance,
                  poisson_model$df.null - poisson_model$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 0.001567133
```

Anova

We can also determine the model terms that cause a significant reduction in deviance.

```
anova(poisson_model, test="Chisq")
```

```
## Analysis of Deviance Table
##
## Model: poisson, link: log
##
## Response: art
##
## Terms added sequentially (first to last)
##
##
##      Df Deviance Resid. Df Resid. Dev Pr(>Chi)
## NULL          914      1817.4
## kid5   3   11.3045      911      1806.1 0.01019 *
## mar    1    6.1638      910      1799.9 0.01304 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Possion Regression with continuous variables, mentor articles and martial status

To model the martial status and number of mentor articles as covariates, I use the formula `art ~ 1+mar + ment`. This formula is equivalent to

$$\log \mu_i = \beta_0 + \beta_1 X_{1i} + X_{2i}$$

where

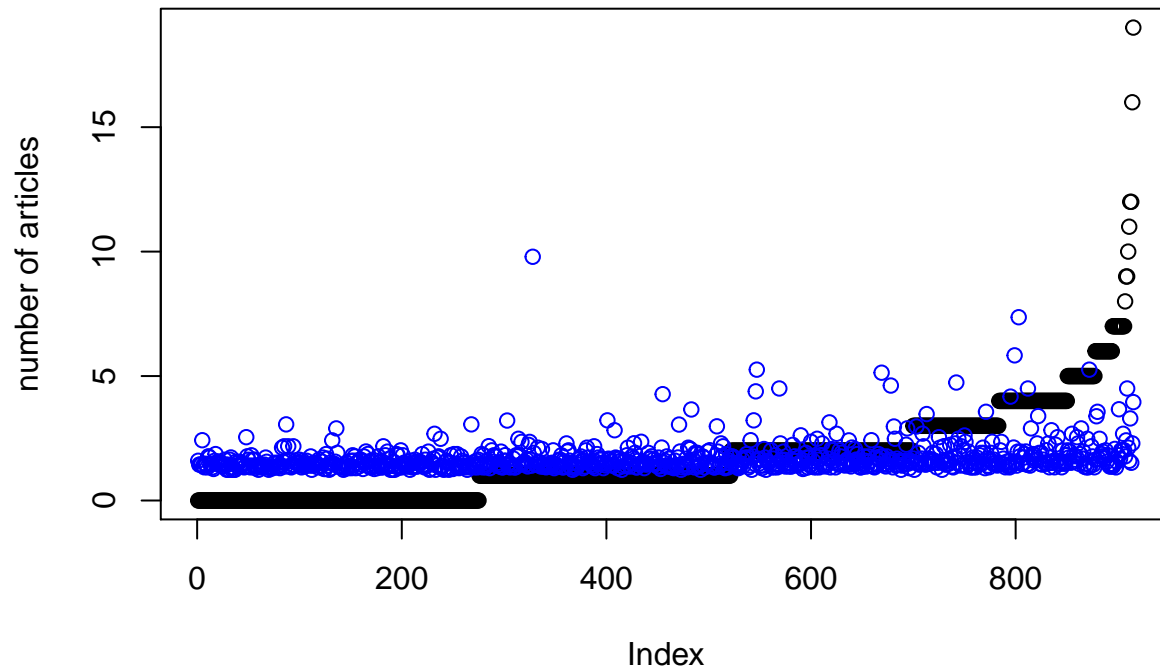
$$X_{1i} = \begin{cases} 1 & \text{if the } i \text{ data point is Married} \\ 0 & \text{otherwise} \end{cases}$$

and X_{2i} is the number of publications of the i th data point's mentor.

```
poisson_model = glm(art ~ 1 + ment + mar,
                    family=poisson(link=log), data=bioChemists)
summary(poisson_model)
```

```
##
## Call:
## glm(formula = art ~ 1 + ment + mar, family = poisson(link = log),
##      data = bioChemists)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.6086  -1.6317  -0.3608   0.5039   5.8942
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.210726   0.049847   4.227 2.36e-05 ***
## ment         0.025917   0.001915  13.530 < 2e-16 ***
## marMarried   0.075332   0.054643   1.379  0.168
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 1817.4  on 914  degrees of freedom
## Residual deviance: 1667.6  on 912  degrees of freedom
```

```
## AIC: 3341.4
##
## Number of Fisher Scoring iterations: 5
plot(bioChemists$art,ylab='number of articles',xlab = 'Index')
points(poisson_model$fitted,col="blue",pch=1)
```



Graphically, we can see that marital status and number of children is not good indicator of number articles published.

Goodness of Fit

Saturated model

We can again compare the current model to the saturated model (best possible fit).

```
p_value = pchisq(poisson_model$deviance,
                 poisson_model$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 6.132629e-47
```

Since our p value is less than 0.05, we reject the null hypothesis. The models are not equivalent.

Null model

We can also compare the current model to the null model (worst possible fit).

```
p_value = pchisq(poisson_model$null.deviance - poisson_model$deviance,
                 poisson_model$df.null - poisson_model$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 2.993003e-33
```

Anova

We can also determine the model terms that cause a significance reduction in deviance.

```
anova(poisson_model, test="Chisq")
```

```
## Analysis of Deviance Table
##
## Model: poisson, link: log
##
## Response: art
##
## Terms added sequentially (first to last)
##
##      Df Deviance Resid. Df Resid. Dev Pr(>Chi)
## NULL                914      1817.4
## ment  1  147.860      913      1669.5 <2e-16 ***
## mar   1    1.918      912      1667.6  0.1661
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Log-Linear Regression

Log-linear models allow us to model association between two or more variables in contingency table. In a log-linear model, there are no well defined explanatory/response variables. This is because we are focused more on the *interaction* between two variables.

Contingency Table

Contingency table displays number of observations for a given combination of factors.

This definition is best represented by an example.

```
bioChemists$art_binary <- sapply(bioChemists$art,function(x) ifelse(x > 1, 1, 0))
bioChemists$ment_binary <- sapply(bioChemists$ment,function(x) ifelse(x > median(bioChemists$ment), 1, 0))
```

One-Way Contingency Table

A one-way contingency table shows the counts according to one covariate.

```
table(art_relative=bioChemists$art_binary)
```

```
## art_relative
##    0    1
## 521 394
```

This one-way contingency table shows that:

- there are 521 biochemists with 1 or less papers
- there are 394 biochemists with greater than 1 papers.

Two-Way Contingency Table

A two-way contingency table shows the counts according to two covariates.

```
table(art_relative=bioChemists$art_binary,ment=bioChemists$ment_binary)
```

```
##           ment
## art_relative  0  1
##           0 321 200
##           1 171 223
```

This two-way contingency table shows that:

- there are 321 biochemists with 1 or less papers and with a mentor that produced less than or equal to 6 papers
- there are 200 biochemists with 1 or less papers and with a mentor that produced more than 6 papers
- there are 171 biochemists with more than 1 paper and with a mentor that produced less than or equal to 6 papers
- there are 223 biochemists with more than 1 paper and with a mentor that produced more than 6 papers

Three-Way Contingency Table

A three-way contingency table shows the counts according to three covariates.

```
table(art_relative=bioChemists$art_binary,ment=bioChemists$ment_binary,
      kid5=bioChemists$kid5)
```

```
## , , kid5 = 0
##
##          ment
## art_relative  0   1
##              0 208 128
##              1 116 147
##
## , , kid5 = 1
##
##          ment
## art_relative  0   1
##              0  66  46
##              1  38  45
##
## , , kid5 = 2
##
##          ment
## art_relative  0   1
##              0  39  21
##              1  16  29
##
## , , kid5 = 3
##
##          ment
## art_relative  0   1
##              0   8   5
##              1   1   2
```

This two-way contingency table shows that:

- With no children,
 - there are 208 biochemists with 1 or less papers and with a mentor that produced less than or equal to 6 papers
 - there are 128 biochemists with 1 or less papers and with a mentor that produced more than 6 papers
 - there are 116 biochemists with more than 1 paper and with a mentor that produced less than or equal to 6 papers
 - there are 147 biochemists with more than 1 paper and with a mentor that produced more than 6 papers
- With 1 child,
 - there are 66 biochemists with 1 or less papers and with a mentor that produced less than or equal to 6 papers
 - there are 46 biochemists with 1 or less papers and with a mentor that produced more than 6 papers
 - there are 38 biochemists with more than 1 paper and with a mentor that produced less than or equal to 6 papers
 - there are 45 biochemists with more than 1 paper and with a mentor that produced more than 6 papers
- With 2 children,
 - there are 39 biochemists with 1 or less papers and with a mentor that produced less than or equal to 6 papers
 - there are 21 biochemists with 1 or less papers and with a mentor that produced more than 6

- papers
 - there are 16 biochemists with more than 1 paper and with a mentor that produced less than or equal to 6 papers
 - there are 29 biochemists with more than 1 paper and with a mentor that produced more than 6 papers
- With 3 children,
 - there are 8 biochemists with 1 or less papers and with a mentor that produced less than or equal to 6 papers
 - there are 5 biochemists with 1 or less papers and with a mentor that produced more than 6 papers
 - there are 1 biochemists with more than 1 paper and with a mentor that produced less than or equal to 6 papers
 - there are 2 biochemists with more than 1 paper and with a mentor that produced more than 6 papers

Log-linear regression

For a two way contingency table, log-linear GLMs have the following properties:

- count response variables, Y_{ij} , which is the number of entries in the (i,j)th cell of the table. Y_{ij} follows a Poisson distribution with mean μ_{ij} .
- a systematic component, η_i , that relates the explanatory variables,

$$\eta_{ij} = \sum_{j=1}^n \beta_k X_{ijk}$$

- a link function that relates the mean of the random to the systematic component

$$\log \mu_{ij} = \sum_{k=1}^n \beta_k X_{ijk}$$

Independent Model for two-way contingency table

We use log-linear model to model the group mean count of each cell of the contingency table. Remember, using a log-linear model, our primary goal is to learn the interaction effects between covariates.

Again, we build the same two-way contingency table. We need to convert the contingency table in a form that is acceptable to glm. We do this below.

```
contingency_table = table(art_relative=bioChemists$art_binary,ment=bioChemists$ment_binary)
contingency_table.df = as.data.frame(contingency_table)
print(contingency_table.df)
```

```
##   art_relative ment Freq
## 1           0    0  321
## 2           1    0  171
## 3           0    1  200
## 4           1    1  223
```

Assuming each number of articles and mentor do not affect each other, we build a model of the cell count that does not take into account interaction effects. Such a model is called the *independent* model. To do this, we use formula `Freq ~ art_relative + ment`.

```
log_linear_model_int <- glm(Freq ~ art_relative + ment,
                           data = contingency_table.df, family = poisson)
summary(log_linear_model_int)
```

```
##
## Call:
## glm(formula = Freq ~ art_relative + ment, family = poisson, data = contingency_table.df)
##
## Deviance Residuals:
##      1      2      3      4
##  2.385 -2.905 -2.713  2.923
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   5.63530    0.05347 105.392 < 2e-16 ***
## art_relative1 -0.27940    0.06676  -4.185 2.85e-05 ***
## ment1         -0.15111    0.06631  -2.279  0.0227 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 52.927  on 3  degrees of freedom
## Residual deviance: 30.035  on 1  degrees of freedom
## AIC: 65.008
##
## Number of Fisher Scoring iterations: 4
```

Goodness of fit

We compare the current model to the saturated model (best possible fit).

```
p_value = pchisq(log_linear_model_int$deviance,
                 log_linear_model_int$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 4.243721e-08
```

Since our p value is less than 0.05, we reject the null hypothesis. The models are not equivalent.

Saturated Model for the two-way contingency table

Assuming each number of articles and mentor affect each other, we build a model of the cell count that takes into account all interaction effects. Such a model is called the *saturated* model. To do this, we use formula `Freq ~ art_relative*ment`.

```
log_linear_model_sat <- glm(Freq ~ art_relative*ment,
                           data = contingency_table.df, family = poisson)
summary(log_linear_model_sat)
```

```
##
## Call:
## glm(formula = Freq ~ art_relative * ment, family = poisson, data = contingency_table.df)
##
## Deviance Residuals:
## [1]  0  0  0  0
##
## Coefficients:
```



```
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)      5.77144    0.05581 103.404 < 2e-16 ***
## art_relative1    -0.62978    0.09467  -6.652 2.89e-11 ***
## ment1           -0.47312    0.09008  -5.252 1.50e-07 ***
## art_relative1:ment1 0.73863    0.13582   5.438 5.38e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
## Null deviance: 5.2927e+01 on 3 degrees of freedom
## Residual deviance: 1.8874e-14 on 0 degrees of freedom
## AIC: 36.973
##
## Number of Fisher Scoring iterations: 2
```

Goodness of fit

We compare the current model to the saturated model (best possible fit).

```
p_value = pchisq(0,
                 log_linear_model_sat$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 1
```

We fail to reject the null hypothesis. This makes sense since the models are literally the same thing.

Model Comparison

We use `anova` with `test='Chisq'` to compare the independent and saturated model.

```
anova(log_linear_model_int, log_linear_model_sat, test='Chisq')
```

```
## Analysis of Deviance Table
##
## Model 1: Freq ~ art_relative + ment
## Model 2: Freq ~ art_relative * ment
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1         1    30.035
## 2         0     0.000  1    30.035 4.244e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

From `anova`, we can see that the saturated model provides a statistically significant result.

We use also `anova` to determine what caused the significant decrease in the deviance.

```
anova(log_linear_model_sat, test='Chisq')
```

```
## Analysis of Deviance Table
##
## Model: poisson, link: log
##
## Response: Freq
##
```

```
## Terms added sequentially (first to last)
##
##
##           Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                        3      52.927
## art_relative      1  17.6844      2    35.243 2.608e-05 ***
## ment              1   5.2082      1    30.035  0.02248 *
## art_relative:ment  1  30.0348      0     0.000 4.244e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Adding `art_relative:ment` to the independent model caused significant decrease in deviance.

Independent Model for the three-way contingency table

Again, we build the same three-way contingency table. We need to convert the contingency table in a form that is acceptable to `glm`. We do this below.

```
contingency_table = table(art_relative=bioChemists$art_binary,
                          ment=bioChemists$ment_binary,
                          kid5=bioChemists$kid5)
contingency_table.df = as.data.frame(contingency_table)
```

To create the *independent* model for the three-way contingency table, we use formula `Freq ~ art_relative + ment + kid5`.

```
log_linear_model_int <- glm(Freq ~ art_relative + ment + kid5,
                           data = contingency_table.df, family = poisson)
summary(log_linear_model_int)
```

```
##
## Call:
## glm(formula = Freq ~ art_relative + ment + kid5, family = poisson,
##      data = contingency_table.df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4439  -1.4070  -0.1702   1.1974   2.4521
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   5.21164    0.05861  88.915 < 2e-16 ***
## art_relative1 -0.27940    0.06676  -4.185 2.85e-05 ***
## ment1         -0.15111    0.06631  -2.279  0.0227 *
## kid51         -1.12226    0.08245 -13.612 < 2e-16 ***
## kid52         -1.74130    0.10580 -16.459 < 2e-16 ***
## kid53         -3.62267    0.25331 -14.301 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 901.879  on 15  degrees of freedom
## Residual deviance:  36.651  on 10  degrees of freedom
## AIC: 131.03
```

```
##
## Number of Fisher Scoring iterations: 4
```

Goodness of fit

We compare the current model to the saturated model (best possible fit).

```
p_value = pchisq(log_linear_model_int$deviance,
                  log_linear_model_int$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 6.50121e-05
```

Since our p value is less than 0.05, we reject the null hypothesis. The models are not equivalent.

Saturated Model

To create the *independent* model for the three-way contingency table, we use formula `Freq ~ art_relative*ment*kid5`.

```
log_linear_model_sat <- glm(Freq ~ art_relative*ment*kid5,
                             data = contingency_table.df, family = poisson)
summary(log_linear_model_sat)
```

```
##
## Call:
## glm(formula = Freq ~ art_relative * ment * kid5, family = poisson,
##      data = contingency_table.df)
##
## Deviance Residuals:
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      5.33754    0.06934  76.979 < 2e-16 ***
## art_relative1    -0.58395    0.11588  -5.039 4.67e-07 ***
## ment1            -0.48551    0.11234  -4.322 1.55e-05 ***
## kid51            -1.14788    0.14128  -8.125 4.47e-16 ***
## kid52            -1.67398    0.17450  -9.593 < 2e-16 ***
## kid53            -3.25810    0.36029  -9.043 < 2e-16 ***
## art_relative1:ment1  0.72235    0.16746   4.314 1.61e-05 ***
## art_relative1:kid51  0.03188    0.23430   0.136  0.892
## art_relative1:kid52 -0.30703    0.31870  -0.963  0.335
## art_relative1:kid53 -1.49549    1.06697  -1.402  0.161
## ment1:kid51        0.12449    0.22251   0.559  0.576
## ment1:kid52       -0.13353    0.29305  -0.456  0.649
## ment1:kid53        0.01550    0.58105   0.027  0.979
## art_relative1:ment1:kid51 -0.19226    0.33686  -0.571  0.568
## art_relative1:ment1:kid52  0.49140    0.44529   1.104  0.270
## art_relative1:ment1:kid53  0.44080    1.36126   0.324  0.746
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
```

```
##
## Null deviance: 9.0188e+02 on 15 degrees of freedom
## Residual deviance: -3.6859e-14 on 0 degrees of freedom
## AIC: 114.38
##
## Number of Fisher Scoring iterations: 3
```

Goodness of fit

We compare the current model to the saturated model (best possible fit).

```
p_value = pchisq(log_linear_model_sat$deviance,
                 log_linear_model_sat$df.residual, lower.tail = F)
print(p_value)
```

```
## [1] 1
```

We fail to reject the null hypothesis. This makes sense since the models are literally the same thing.

Model Comparison

We use `anova` with `test='Chisq'` to compare the independent and saturated model.

```
anova(log_linear_model_int, log_linear_model_sat, test='Chisq')
```

```
## Analysis of Deviance Table
##
## Model 1: Freq ~ art_relative + ment + kid5
## Model 2: Freq ~ art_relative * ment * kid5
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1         10      36.651
## 2          0       0.000 10   36.651 6.501e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

From `anova`, we can see that the saturated model provides a statistically significant result.

We use also `anova` to determine what caused the significant decrease in the deviance.

```
anova(log_linear_model_sat, test='Chisq')
```

```
## Analysis of Deviance Table
##
## Model: poisson, link: log
##
## Response: Freq
##
## Terms added sequentially (first to last)
##
##
##           Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                        15      901.88
## art_relative          1      17.68      14      884.20 2.608e-05 ***
## ment                  1       5.21      13      878.99 0.02248 *
## kid5                   3      842.34      10      36.65 < 2.2e-16 ***
## art_relative:ment      1       30.03       9       6.62 4.244e-08 ***
```

```
## art_relative:kid5      3      4.45      6      2.17  0.21665
## ment:kid5              3      0.19      3      1.97  0.97873
## art_relative:ment:kid5 3      1.97      0      0.00  0.57819
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Adding `art_relative:ment` to the independent model caused significant decrease in deviance.

Hierarchical modeling


```
student_data <- read.csv("data/hsb1.csv")
school_data <- read.csv("data/hsb2.csv")
student_data$ses_grandmean <- student_data$ses - mean(student_data$ses) # Grand-mean centered student SES
school_data$sm_ses_grandmean <- school_data$meanses - mean(school_data$meanses) # Grand-mean centered school SES

data <- merge(student_data, school_data, by = "id")

ses_group_mean <- aggregate(data$ses, list(data$id), FUN = mean, data = data) # Group-mean centered student SES
names(ses_group_mean) <- c('id', 'groupmeanSES')
data <- merge(data, ses_group_mean, by = "id")

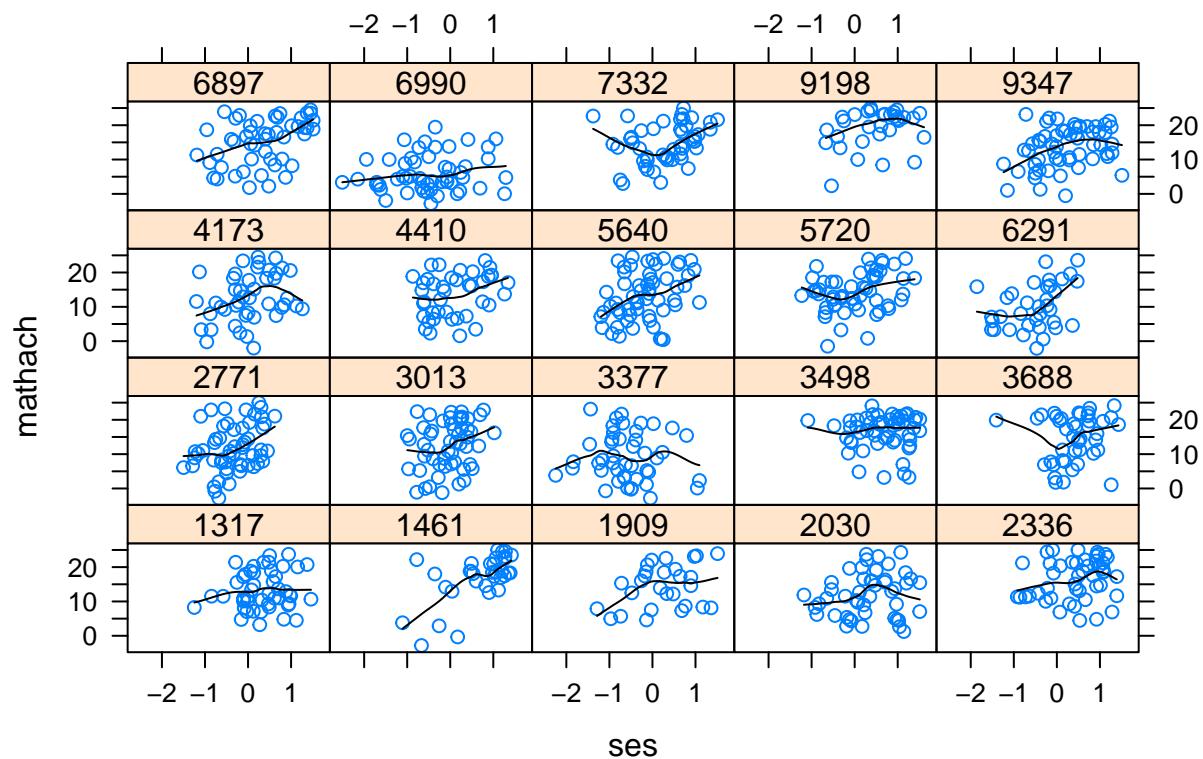
groups <- unique(data$id)[sample(1:160, 20)]
subset <- data[data$id %in% groups, ]

xyplot(mathach ~ ses | as.factor(id), subset,
       col.line = 'black',
       type = c("p", "r"),
       main = 'Variability in Math Achievement ~ SES Relationship')
```




```
xyplot(mathach ~ ses | as.factor(id), subset,
       col.line = 'black',
       type = c("p", "smooth"),
       main = 'Variability in Math Achievement ~ SES Relationship')
```

Variability in Math Achievement ~ SES Relationship



```
xyplot(mathach ~ ses, subset,
       type = c("p", "smooth"),
       group = data$id,
       main = 'Variability in Math Achievement ~ SES Relationship')
```

Variability in Math Achievement ~ SES Relationship



```
unconditional <- lmer(mathach ~ 1 + (1|id), data = data)
summary(unconditional) # on p-values in nlme: https://stat.ethz.ch/pipermail/r-help/2006-May/094765.htm

## Linear mixed model fit by REML ['lmerMod']
## Formula: mathach ~ 1 + (1 | id)
## Data: data
##
## REML criterion at convergence: 47116.8
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.0631 -0.7539  0.0267  0.7606  2.7426
##
## Random effects:
## Groups Name Variance Std.Dev.
## id      (Intercept) 8.614 2.935
## Residual          39.148 6.257
## Number of obs: 7185, groups: id, 160
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept) 12.6370    0.2444  51.71

confint(unconditional) # you can also just calculate an approximate 95% confidence interval yourself: e

## Computing profile confidence intervals ...

##              2.5 %    97.5 %
## .sig01        2.594729 3.315880
```



```

## .sigma      6.154803  6.361786
## (Intercept) 12.156289 13.117121

unconditional_2 <- lme(mathach ~ 1, random = ~ 1 | id, data = data)
summary(unconditional_2)

## Linear mixed-effects model fit by REML
## Data: data
##      AIC      BIC    logLik
##  47122.79 47143.43 -23558.4
##
## Random effects:
## Formula: ~1 | id
##      (Intercept) Residual
## StdDev:      2.934966 6.256862
##
## Fixed effects: mathach ~ 1
##              Value Std.Error   DF  t-value p-value
## (Intercept) 12.63697 0.2443936 7025 51.70747      0
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.06312473 -0.75387398  0.02670132  0.76062171  2.74262579
##
## Number of Observations: 7185
## Number of Groups: 160

random_intercept_fixed_slope <- lmer(mathach ~ 1 + groupmeanSES + (1|id), data = data)
summary(random_intercept_fixed_slope)

## Linear mixed model fit by REML ['lmerMod']
## Formula: mathach ~ 1 + groupmeanSES + (1 | id)
## Data: data
##
## REML criterion at convergence: 46961.3
##
## Scaled residuals:
##      Min      1Q   Median      3Q      Max
## -3.13493 -0.75254  0.02413  0.76766  2.78515
##
## Random effects:
## Groups   Name      Variance Std.Dev.
## id      (Intercept)  2.639    1.624
## Residual                39.157    6.258
## Number of obs: 7185, groups: id, 160
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)   12.6846    0.1493   84.97
## groupmeanSES    5.8635    0.3615   16.22
##
## Correlation of Fixed Effects:
##              (Intr)
## groupmenSES 0.010

```

```

confint(random_intercept_fixed_slope)

## Computing profile confidence intervals ...
##           2.5 %    97.5 %
## .sig01      1.385193  1.871127
## .sigma      6.155502  6.362511
## (Intercept) 12.391774 12.976903
## groupmeanSES 5.155743 6.572440

random_intercept_fixed_slope_2 <- lme(mathach ~ 1 + groupmeanSES, random = ~ 1 | id, data = data)
summary(random_intercept_fixed_slope_2)

## Linear mixed-effects model fit by REML
## Data: data
##      AIC      BIC    logLik
## 46969.29 46996.81 -23480.65
##
## Random effects:
## Formula: ~1 | id
##      (Intercept) Residual
## StdDev:      1.624462 6.257562
##
## Fixed effects: mathach ~ 1 + groupmeanSES
##              Value Std.Error   DF  t-value p-value
## (Intercept) 12.684609 0.1492900 7025 84.96624      0
## groupmeanSES 5.863539 0.3614712  158 16.22132      0
## Correlation:
##      (Intr)
## groupmeanSES 0.01
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.13493066 -0.75254260  0.02413095  0.76766113  2.78515398
##
## Number of Observations: 7185
## Number of Groups: 160

random_intercept_random_slope <- lmer(mathach ~ 1 + groupmeanSES + (1 + groupmeanSES|id), data = data)

## boundary (singular) fit: see ?isSingular

summary(random_intercept_random_slope)

## Linear mixed model fit by REML ['lmerMod']
## Formula: mathach ~ 1 + groupmeanSES + (1 + groupmeanSES | id)
## Data: data
##
## REML criterion at convergence: 46960.9
##
## Scaled residuals:
##      Min      1Q  Median      3Q      Max
## -3.13245 -0.75164  0.02212  0.76876  2.79449
##
## Random effects:
## Groups   Name              Variance Std.Dev. Corr

```

```

## id      (Intercept)  2.62707 1.6208
##      groupmeanSES  0.05417 0.2327  -1.00
## Residual              39.15798 6.2576
## Number of obs: 7185, groups: id, 160
##
## Fixed effects:
##      Estimate Std. Error t value
## (Intercept)  12.6832    0.1491  85.04
## groupmeanSES  5.8379    0.3644  16.02
##
## Correlation of Fixed Effects:
##      (Intr)
## groupmenSES -0.078
## convergence code: 0
## boundary (singular) fit: see ?isSingular
random_intercept_random_slope_2 <- lme(mathach ~ 1 + groupmeanSES, random = ~ 1 + groupmeanSES | id, da
summary(random_intercept_random_slope_2)

## Linear mixed-effects model fit by REML
## Data: data
##      AIC      BIC    logLik
## 46973.29 47014.57 -23480.65
##
## Random effects:
## Formula: ~1 + groupmeanSES | id
## Structure: General positive-definite, Log-Cholesky parametrization
##      StdDev      Corr
## (Intercept) 1.624460932 (Intr)
## groupmeanSES 0.008272356 -0.003
## Residual    6.257561467
##
## Fixed effects: mathach ~ 1 + groupmeanSES
##      Value Std.Error DF t-value p-value
## (Intercept) 12.684610 0.1492901 7025 84.96616 0
## groupmeanSES 5.863533 0.3614729 158 16.22122 0
## Correlation:
##      (Intr)
## groupmeanSES 0.01
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.13493049 -0.75254293 0.02413128 0.76766157 2.78515572
##
## Number of Observations: 7185
## Number of Groups: 160
fixed_intercept_random_slope <- lmer(mathach ~ 1 + groupmeanSES + (0 + groupmeanSES|id), data = data)
summary(fixed_intercept_random_slope)

## Linear mixed model fit by REML ['lmerMod']
## Formula: mathach ~ 1 + groupmeanSES + (0 + groupmeanSES | id)
## Data: data
##
## REML criterion at convergence: 47065

```

```

##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.1150 -0.7431  0.0317  0.7651  2.8202
##
## Random effects:
##   Groups   Name                Variance Std.Dev.
##    id      groupmeanSES 27.05      5.201
## Residual                39.75      6.304
## Number of obs: 7185, groups: id, 160
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  12.7640    0.1226  104.07
## groupmeanSES   5.4202    0.5271   10.28
##
## Correlation of Fixed Effects:
##              (Intr)
## groupmenSES -0.045
fixed_intercept_random_slope_2 <- lme(mathach ~ 1 + groupmeanSES, random = ~ 0 + groupmeanSES | id, data = data,
summary(fixed_intercept_random_slope_2)

## Linear mixed-effects model fit by REML
## Data: data
##      AIC      BIC    logLik
## 47072.99 47100.51 -23532.5
##
## Random effects:
## Formula: ~0 + groupmeanSES | id
##      groupmeanSES Residual
## StdDev:      5.201045 6.304462
##
## Fixed effects: mathach ~ 1 + groupmeanSES
##              Value Std.Error   DF   t-value p-value
## (Intercept) 12.764014 0.1226493 7025 104.06918      0
## groupmeanSES  5.420157 0.5270957 158 10.28306      0
## Correlation:
##              (Intr)
## groupmeanSES -0.045
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.11504273 -0.74308714  0.03169931  0.76511017  2.82021818
##
## Number of Observations: 7185
## Number of Groups: 160
fixed_slope_level_two_variable <- lmer(mathach ~ 1 + groupmeanSES + sm_ses_grandmean + (1|id), data = data,
summary(fixed_slope_level_two_variable)

## Linear mixed model fit by REML ['lmerMod']
## Formula: mathach ~ 1 + groupmeanSES + sm_ses_grandmean + (1 | id)
## Data: data
##

```

```

## REML criterion at convergence: 46946.8
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.13127 -0.75215  0.02439  0.76700  2.78177
##
## Random effects:
##   Groups   Name            Variance Std.Dev.
##   id       (Intercept)    2.659    1.631
##   Residual                39.157    6.258
## Number of obs: 7185, groups: id, 160
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)      11.675      3.299   3.539
## groupmeanSES     -157.361    532.675  -0.295
## sm_ses_grandmean  163.223    532.668   0.306
##
## Correlation of Fixed Effects:
##              (Intr) grpSES
## groupmenSES   0.999
## sm_ss_grndm  -0.999 -1.000
fixed_slope_level_two_variable_2 <- lme(mathach ~ 1 + groupmeanSES + sm_ses_grandmean, random = ~ 1 | id, data = data)
summary(fixed_slope_level_two_variable_2)

## Linear mixed-effects model fit by REML
## Data: data
##      AIC      BIC    logLik
## 46956.81 46991.2 -23473.4
##
## Random effects:
## Formula: ~1 | id
##      (Intercept) Residual
## StdDev:      1.630771 6.257562
##
## Fixed effects: mathach ~ 1 + groupmeanSES + sm_ses_grandmean
##              Value Std.Error   DF   t-value p-value
## (Intercept)      11.67469    3.2988 7025   3.539111  0.0004
## groupmeanSES     -157.36077   532.6748  157  -0.295416  0.7681
## sm_ses_grandmean  163.22262   532.6683  157   0.306425  0.7597
## Correlation:
##              (Intr) grpSES
## groupmeanSES      0.999
## sm_ses_grandmean -0.999 -1.000
##
## Standardized Within-Group Residuals:
##      Min       Q1       Med       Q3      Max
## -3.13126623 -0.75215319  0.02439264  0.76699775  2.78176653
##
## Number of Observations: 7185
## Number of Groups: 160

```

```

random_slope_level_two_variable <- lmer(mathach ~ 1 + groupmeanSES + sm_ses_grandmean + (1 + groupmeanSES |
## boundary (singular) fit: see ?isSingular
summary(random_slope_level_two_variable)

## Linear mixed model fit by REML ['lmerMod']
## Formula:
## mathach ~ 1 + groupmeanSES + sm_ses_grandmean + (1 + groupmeanSES |
## id)
## Data: data
##
## REML criterion at convergence: 46946.3
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.12727 -0.74930  0.02286  0.76841  2.79122
##
## Random effects:
##   Groups   Name                Variance Std.Dev. Corr
##   id       (Intercept)         2.64688  1.6269
##           groupmeanSES         0.05901  0.2429   -1.00
##   Residual                        39.15801  6.2576
## Number of obs: 7185, groups: id, 160
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)      11.493      3.292   3.491
## groupmeanSES     -186.500     531.608  -0.351
## sm_ses_grandmean  192.333     531.597   0.362
##
## Correlation of Fixed Effects:
##              (Intr) grpSES
## groupmenSES   0.999
## sm_ss_grndm  -0.999 -1.000
## convergence code: 0
## boundary (singular) fit: see ?isSingular
random_slope_level_two_variable_2 <- lme(mathach ~ 1 + groupmeanSES + sm_ses_grandmean, random = ~ 1 +
summary(random_slope_level_two_variable_2)

## Linear mixed-effects model fit by REML
## Data: data
##      AIC      BIC    logLik
## 46960.81 47008.96 -23473.4
##
## Random effects:
## Formula: ~1 + groupmeanSES | id
## Structure: General positive-definite, Log-Cholesky parametrization
##              StdDev    Corr
## (Intercept)  1.6307657 (Intr)
## groupmeanSES 0.0130297 -0.005
## Residual     6.2575620
##
## Fixed effects: mathach ~ 1 + groupmeanSES + sm_ses_grandmean

```

```

##              Value Std.Error   DF   t-value p-value
## (Intercept)    11.67466    3.2988 7025  3.539092  0.0004
## groupmeanSES   -157.36607   532.6762  157 -0.295425  0.7681
## sm_ses_grandmean 163.22790   532.6697  157  0.306434  0.7597
## Correlation:
##              (Intr) grpSES
## groupmeanSES    0.999
## sm_ses_grandmean -0.999 -1.000
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.13126549 -0.75215462  0.02439169  0.76699807  2.78177069
##
## Number of Observations: 7185
## Number of Groups: 160

fixed_slope_cl_interaction <- lmer(mathach ~ 1 + groupmeanSES*sm_ses_grandmean + (1|id), data = data)
summary(fixed_slope_cl_interaction)

## Linear mixed model fit by REML ['lmerMod']
## Formula: mathach ~ 1 + groupmeanSES * sm_ses_grandmean + (1 | id)
## Data: data
##
## REML criterion at convergence: 46945
##
## Scaled residuals:
##      Min      1Q   Median      3Q      Max
## -3.11930 -0.75112  0.02448  0.76597  2.78831
##
## Random effects:
## Groups Name Variance Std.Dev.
## id      (Intercept) 2.664  1.632
## Residual          39.158  6.258
## Number of obs: 7185, groups: id, 160
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)      11.4252    3.3155  3.446
## groupmeanSES     -213.6963   537.6427 -0.397
## sm_ses_grandmean  219.4884   537.6248  0.408
## groupmeanSES:sm_ses_grandmean -0.5799    0.7253 -0.800
##
## Correlation of Fixed Effects:
##              (Intr) grpSES sm_ss_
## groupmenSES  0.998
## sm_ss_grndm -0.998 -1.000
## grpmnSES:__  0.094  0.131 -0.131

fixed_slope_cl_interaction_2 <- lme(mathach ~ 1 + groupmeanSES*sm_ses_grandmean, random = ~ 1 | id, data = data)
summary(fixed_slope_cl_interaction_2)

## Linear mixed-effects model fit by REML
## Data: data
##      AIC      BIC    logLik
## 46956.97 46998.25 -23472.49

```

```
##
## Random effects:
## Formula: ~1 | id
## (Intercept) Residual
## StdDev: 1.632105 6.257638
##
## Fixed effects: mathach ~ 1 + groupmeanSES * sm_ses_grandmean
##
```

	Value	Std.Error	DF	t-value	p-value
(Intercept)	11.42519	3.3155	7025	3.445968	0.0006
groupmeanSES	-213.69625	537.6427	156	-0.397469	0.6916
sm_ses_grandmean	219.48842	537.6248	156	0.408256	0.6836
groupmeanSES:sm_ses_grandmean	-0.57991	0.7253	156	-0.799543	0.4252

```
## Correlation:
##
```

	(Intr)	grpSES	sm_ss_
groupmeanSES	0.998		
sm_ses_grandmean	-0.998	-1.000	
groupmeanSES:sm_ses_grandmean	0.094	0.131	-0.131

```
##
## Standardized Within-Group Residuals:
## Min Q1 Med Q3 Max
## -3.11929841 -0.75112002 0.02448373 0.76596673 2.78831371
##
## Number of Observations: 7185
## Number of Groups: 160

random_slope_cl_interaction <- lmer(mathach ~ 1 + groupmeanSES*sm_ses_grandmean + (1 + groupmeanSES|id)

## boundary (singular) fit: see ?isSingular

summary(random_slope_cl_interaction)

## Linear mixed model fit by REML ['lmerMod']
## Formula:
## mathach ~ 1 + groupmeanSES * sm_ses_grandmean + (1 + groupmeanSES |
## id)
## Data: data
##
## REML criterion at convergence: 46944.6
##
## Scaled residuals:
## Min 1Q Median 3Q Max
## -3.11654 -0.75065 0.02247 0.76812 2.79659
##
## Random effects:
## Groups Name Variance Std.Dev. Corr
## id (Intercept) 2.65355 1.6290
## groupmeanSES 0.04692 0.2166 -1.00
## Residual 39.15898 6.2577
## Number of obs: 7185, groups: id, 160
##
## Fixed effects:
##
```

	Estimate	Std. Error	t value
(Intercept)	11.3252	3.3070	3.425
groupmeanSES	-228.1145	536.0175	-0.426
sm_ses_grandmean	233.9132	536.0005	0.436


```

## groupmeanSES:sm_ses_grandmean    -0.5251      0.7363   -0.713
##
## Correlation of Fixed Effects:
##          (Intr) grpSES sm_ss_
## groupmenSES    0.998
## sm_ss_grndm   -0.998 -1.000
## grpmnSES:__    0.080  0.118 -0.118
## convergence code: 0
## boundary (singular) fit: see ?isSingular

random_slope_cl_interaction_2 <- lme(mathach ~ 1 + groupmeanSES*sm_ses_grandmean, random = ~ 1 + groupm
summary(random_slope_cl_interaction_2)

## Linear mixed-effects model fit by REML
## Data: data
##      AIC      BIC    logLik
## 46960.97 47016.01 -23472.49
##
## Random effects:
## Formula: ~1 + groupmeanSES | id
## Structure: General positive-definite, Log-Cholesky parametrization
##              StdDev      Corr
## (Intercept)  1.632105137 (Intr)
## groupmeanSES 0.005745282 -0.002
## Residual      6.257637586
##
## Fixed effects: mathach ~ 1 + groupmeanSES * sm_ses_grandmean
##              Value Std.Error   DF   t-value p-value
## (Intercept)    11.42519    3.3155 7025   3.445964 0.0006
## groupmeanSES   -213.69606   537.6432  156  -0.397468 0.6916
## sm_ses_grandmean    219.48823   537.6253  156   0.408255 0.6836
## groupmeanSES:sm_ses_grandmean  -0.57990    0.7253  156  -0.799534 0.4252
## Correlation:
##              (Intr) grpSES sm_ss_
## groupmeanSES      0.998
## sm_ses_grandmean  -0.998 -1.000
## groupmeanSES:sm_ses_grandmean  0.094  0.131 -0.131
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.11929848 -0.75111989  0.02448355  0.76596678  2.78831459
##
## Number of Observations: 7185
## Number of Groups: 160

logit_random_intercept_and_slope <- glmer(minority ~ groupmeanSES + (1 + groupmeanSES | id), data = data,
family = binomial(link="logit"))
summary(logit_random_intercept_and_slope)

## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: minority ~ groupmeanSES + (1 + groupmeanSES | id)
## Data: data
##

```

```

##      AIC      BIC   logLik deviance df.resid
##  5453.9   5488.3  -2721.9   5443.9     7180
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -6.2886 -0.3942 -0.2073  0.1590  6.1544
##
## Random effects:
##   Groups Name      Variance Std.Dev. Corr
##   id      (Intercept)  2.529   1.590
##   groupmeanSES 11.445   3.383   -0.32
## Number of obs: 7185, groups: id, 160
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -1.7382     0.1678 -10.359 < 2e-16 ***
## groupmeanSES  -2.0523     0.5370  -3.822 0.000132 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr)
## groupmenSES -0.230
specified_variance_covariance_matrix_for_random_effects <- lme(mathach ~ 1 + groupmeanSES*sm_ses_grandmean,
                                                                correlation = corAR1(), data = data) # j
summary(specified_variance_covariance_matrix_for_random_effects)

## Linear mixed-effects model fit by REML
## Data: data
##      AIC      BIC   logLik
##  46962.8 47024.71 -23472.4
##
## Random effects:
## Formula: ~1 + groupmeanSES | id
## Structure: General positive-definite, Log-Cholesky parametrization
##              StdDev      Corr
## (Intercept)  1.62876592 (Intr)
## groupmeanSES 0.07080922 -0.039
## Residual     6.25836130
##
## Correlation Structure: AR(1)
## Formula: ~1 | id
## Parameter estimate(s):
##      Phi
## 0.005104377
## Fixed effects: mathach ~ 1 + groupmeanSES * sm_ses_grandmean
##              Value Std.Error DF   t-value p-value
## (Intercept)    11.42446    3.3152 7025   3.446083  0.0006
## groupmeanSES   -213.84493   537.5895  156  -0.397785  0.6913
## sm_ses_grandmean  219.63651   537.5716  156   0.408572  0.6834
## groupmeanSES:sm_ses_grandmean  -0.57970    0.7256  156  -0.798894  0.4256
## Correlation:
##              (Intr) grpSES sm_ss_
## groupmeanSES    0.998

```

```

## sm_ses_grandmean          -0.998 -1.000
## groupmeanSES:sm_ses_grandmean  0.094  0.131 -0.131
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -3.11860390 -0.75163778  0.02452874  0.76588109  2.78744053
##
## Number of Observations: 7185
## Number of Groups: 160

```